

# Skin Cancer (Melanoma) Detection



```
import numpy as np
import pandas as pd

malignant_dir =
'/kaggle/input/melanoma-cancer-dataset/train/Malignant'
benign_dir = '/kaggle/input/melanoma-cancer-dataset/train/Benign'

def load_image_paths_labels(malignant_dir, benign_dir):
    malignant_images = [(os.path.join(malignant_dir, img),
'Malignant') for img in os.listdir(malignant_dir)]
    benign_images = [(os.path.join(benign_dir, img), 'Benign') for img
in os.listdir(benign_dir)]
    images_labels = malignant_images + benign_images
    return images_labels

image_data = load_image_paths_labels(malignant_dir, benign_dir)
df = pd.DataFrame(image_data, columns=['image_path', 'label'])
df.head()
```

```

                                image_path    label
0  /kaggle/input/melanoma-cancer-dataset/train/Ma...  Malignant
1  /kaggle/input/melanoma-cancer-dataset/train/Ma...  Malignant
2  /kaggle/input/melanoma-cancer-dataset/train/Ma...  Malignant
3  /kaggle/input/melanoma-cancer-dataset/train/Ma...  Malignant
4  /kaggle/input/melanoma-cancer-dataset/train/Ma...  Malignant

df.tail()

                                image_path    label
11874  /kaggle/input/melanoma-cancer-dataset/train/Be...  Benign
11875  /kaggle/input/melanoma-cancer-dataset/train/Be...  Benign
11876  /kaggle/input/melanoma-cancer-dataset/train/Be...  Benign
11877  /kaggle/input/melanoma-cancer-dataset/train/Be...  Benign
11878  /kaggle/input/melanoma-cancer-dataset/train/Be...  Benign

df.shape
(11879, 2)

df.columns
Index(['image_path', 'label'], dtype='object')

df.duplicated().sum()
0

df.isnull().sum()
image_path    0
label         0
dtype: int64

df['label'].unique()
array(['Malignant', 'Benign'], dtype=object)

df['label'].value_counts()

label
Benign      6289
Malignant   5590
Name: count, dtype: int64

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(15, 8))

ax = sns.countplot(x='label', data=df, palette='hls')
ax.set_ylim(0, df['label'].value_counts().max() + 150)


```

```

for p in ax.patches:
    height = p.get_height()
    ax.annotate(f'{height}',
                xy=(p.get_x() + p.get_width() / 2., height),
                xytext=(0, 10),
                textcoords='offset points', ha='center', va='center')

plt.show()

```



```

label_counts = df['label'].value_counts()


plt.figure(figsize=(8, 8))

plt.pie(label_counts,
        labels=label_counts.index,
        autopct='%1.1f%%',
        startangle=90,
        colors=['#ff9999', '#66b3ff'],
        shadow=True)

plt.axis('equal')
plt.title('Distribution of Malignant and Benign Melanoma Cases')
plt.show()

```

Distribution of Malignant and Benign Melanoma Cases



```
malignant_images = os.listdir(malignant_dir)
benign_images = os.listdir(benign_dir)

malignant_sample = np.random.choice(malignant_images, 5,
replace=False)
benign_sample = np.random.choice(benign_images, 5, replace=False)

import matplotlib.image as mpimg

fig, axes = plt.subplots(2, 5, figsize=(20, 8))


for ax, img_name in zip(axes[0], malignant_sample):
    img_path = os.path.join(malignant_dir, img_name)
    img = mpimg.imread(img_path)
    ax.imshow(img)
    ax.axis('off')
    ax.set_title('Malignant')
```

```

for ax, img_name in zip(axes[1], benign_sample):
    img_path = os.path.join(benign_dir, img_name)
    img = mpimg.imread(img_path)
    ax.imshow(img)
    ax.axis('off')
    ax.set_title('Benign')

plt.tight_layout()
plt.show()

```



```

from imblearn.over_sampling import RandomOverSampler

ros = RandomOverSampler(random_state=42)
X_resampled, y_resampled = ros.fit_resample(df[['image_path']],
df['label'])

df_resampled = pd.DataFrame(X_resampled, columns=['image_path'])
df_resampled['label'] = y_resampled

print("\nClass distribution after oversampling:")
print(df_resampled['label'].value_counts())

```

```

Class distribution after oversampling:
label
Malignant    6289
Benign       6289
Name: count, dtype: int64

```

```

import time
import shutil
import pathlib
import itertools
from PIL import Image

```

```

import cv2
import seaborn as sns
sns.set_style('darkgrid')
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam, Adamax
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense, Activation, Dropout, BatchNormalization
from tensorflow.keras import regularizers

import warnings
warnings.filterwarnings("ignore")

print ('check')

check

train_df_new, temp_df_new = train_test_split(
    df_resampled,
    train_size=0.8,
    shuffle=True,
    random_state=42,
    stratify=df_resampled['label']
)

valid_df_new, test_df_new = train_test_split(
    temp_df_new,
    test_size=0.5,
    shuffle=True,
    random_state=42,
    stratify=temp_df_new['label']
)

from tensorflow.keras.preprocessing.image import ImageDataGenerator

batch_size = 16
img_size = (224, 224)
channels = 3
img_shape = (img_size[0], img_size[1], channels)

tr_gen = ImageDataGenerator(rescale=1./255)
ts_gen = ImageDataGenerator(rescale=1./255)

train_gen_new = tr_gen.flow_from_dataframe(

```

```

    train_df_new,
    x_col='image_path',
    y_col='label',
    target_size=img_size,
    class_mode='binary',
    color_mode='rgb',
    shuffle=True,
    batch_size=batch_size
)

valid_gen_new = ts_gen.flow_from_dataframe(
    valid_df_new,
    x_col='image_path',
    y_col='label',
    target_size=img_size,
    class_mode='binary',
    color_mode='rgb',
    shuffle=True,
    batch_size=batch_size
)

test_gen_new = ts_gen.flow_from_dataframe(
    test_df_new,
    x_col='image_path',
    y_col='label',
    target_size=img_size,
    class_mode='binary',
    color_mode='rgb',
    shuffle=False,
    batch_size=batch_size
)

Found 10062 validated image filenames belonging to 2 classes.
Found 1258 validated image filenames belonging to 2 classes.
Found 1258 validated image filenames belonging to 2 classes.

import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import EarlyStopping

physical_devices = tf.config.list_physical_devices('GPU')
if physical_devices:
    print("Using GPU")
else:
    print("Using CPU")

Using GPU

def create_cnn_model(input_shape):
    model = models.Sequential()

```

```

    model.add(layers.Conv2D(32, (3, 3), activation='relu',
input_shape=input_shape))
    model.add(layers.MaxPooling2D((2, 2)))

    model.add(layers.Conv2D(64, (3, 3), activation='relu'))
    model.add(layers.MaxPooling2D((2, 2)))

    model.add(layers.Conv2D(128, (3, 3), activation='relu'))
    model.add(layers.MaxPooling2D((2, 2)))

    model.add(layers.Flatten())

    model.add(layers.Dense(128, activation='relu'))
    model.add(layers.Dense(1, activation='sigmoid'))

    return model

```

```

input_shape = (224, 224, 3)

cnn_model = create_cnn_model(input_shape)

cnn_model.compile(optimizer='adam',
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

```

```
cnn_model.summary()
```

Model: "sequential"

Layer (type) Param #	Output Shape	
conv2d (Conv2D) 896	(None, 222, 222, 32)	
max_pooling2d (MaxPooling2D) 0	(None, 111, 111, 32)	
conv2d_1 (Conv2D) 18,496	(None, 109, 109, 64)	
max_pooling2d_1 (MaxPooling2D) 0	(None, 54, 54, 64)	



conv2d_2 (Conv2D)	(None, 52, 52, 128)	
73,856		
max_pooling2d_2 (MaxPooling2D)	(None, 26, 26, 128)	
0		
flatten (Flatten)	(None, 86528)	
0		
dense (Dense)	(None, 128)	
11,075,712		
dense_1 (Dense)	(None, 1)	
129		

Total params: 11,169,089 (42.61 MB)

Trainable params: 11,169,089 (42.61 MB)

Non-trainable params: 0 (0.00 B)

```
early_stopping = EarlyStopping(monitor='val_loss', patience=5,
restore_best_weights=True)
```

```
history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=10,
    callbacks=[early_stopping],
    verbose=1
)
```

Epoch 1/10

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

```
I0000 00:00:1729588558.928415    495 service.cc:145] XLA service
0x7e93140049a0 initialized for platform CUDA (this does not guarantee
that XLA will be used). Devices:
I0000 00:00:1729588558.928474    495 service.cc:153]   StreamExecutor
device (0): Tesla T4, Compute Capability 7.5
I0000 00:00:1729588558.928478    495 service.cc:153]   StreamExecutor
device (1): Tesla T4, Compute Capability 7.5
```

```
3/629 _____ 34s 55ms/step - accuracy: 0.5243 - loss: 1.3999
```

```
I0000 00:00:1729588563.166459      495 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.
```

```
629/629 _____ 67s 96ms/step - accuracy: 0.7391 - loss: 0.5452 - val_accuracy: 0.8537 - val_loss: 0.3358
```

```
Epoch 2/10
```

```
629/629 _____ 21s 33ms/step - accuracy: 0.8369 - loss: 0.3690 - val_accuracy: 0.8426 - val_loss: 0.3451
```

```
Epoch 3/10
```

```
629/629 _____ 20s 31ms/step - accuracy: 0.8601 - loss: 0.3302 - val_accuracy: 0.8458 - val_loss: 0.3568
```

```
Epoch 4/10
```

```
629/629 _____ 20s 32ms/step - accuracy: 0.8504 - loss: 0.3340 - val_accuracy: 0.8593 - val_loss: 0.3358
```

```
Epoch 5/10
```

```
629/629 _____ 20s 31ms/step - accuracy: 0.8569 - loss: 0.3128 - val_accuracy: 0.8625 - val_loss: 0.3214
```

```
Epoch 6/10
```

```
629/629 _____ 19s 30ms/step - accuracy: 0.8648 - loss: 0.3169 - val_accuracy: 0.8323 - val_loss: 0.3829
```

```
Epoch 7/10
```

```
629/629 _____ 20s 32ms/step - accuracy: 0.8763 - loss: 0.2824 - val_accuracy: 0.8593 - val_loss: 0.3239
```

```
Epoch 8/10
```

```
629/629 _____ 21s 33ms/step - accuracy: 0.8837 - loss: 0.2724 - val_accuracy: 0.8808 - val_loss: 0.3216
```

```
Epoch 9/10
```

```
629/629 _____ 21s 32ms/step - accuracy: 0.8998 - loss: 0.2386 - val_accuracy: 0.8712 - val_loss: 0.3197
```

```
Epoch 10/10
```

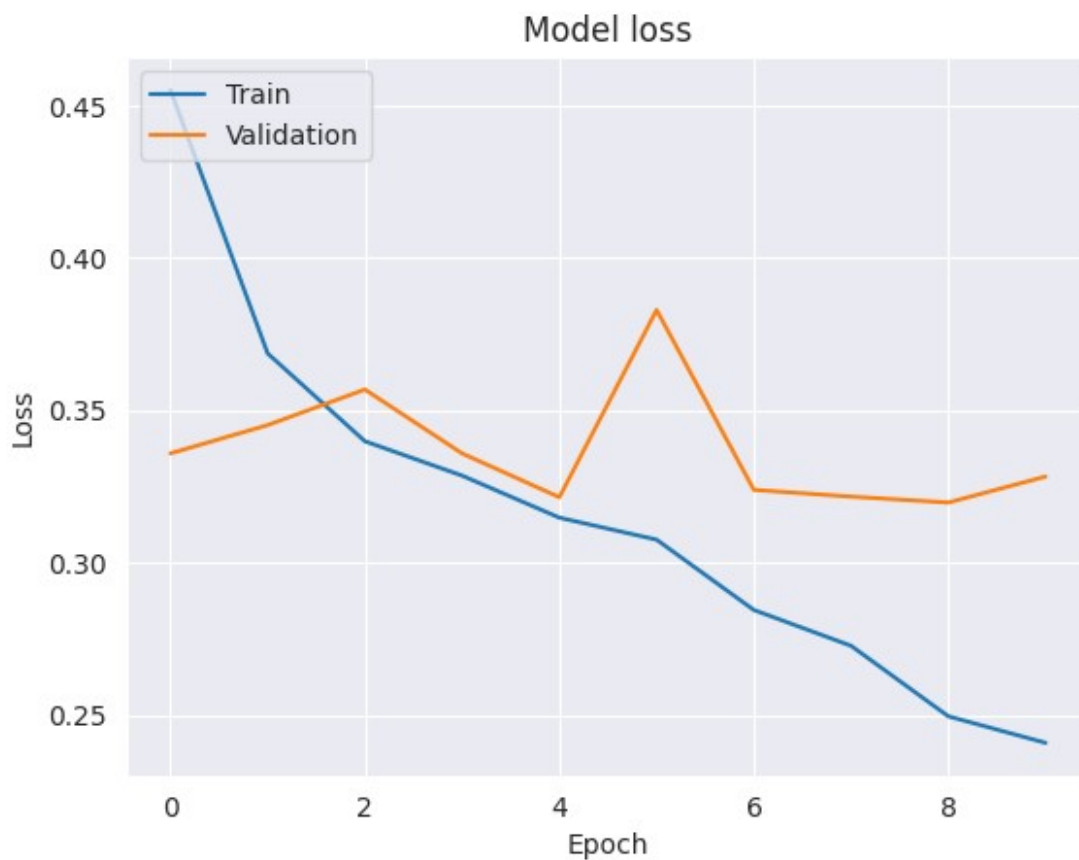
```
629/629 _____ 20s 32ms/step - accuracy: 0.9007 - loss: 0.2340 - val_accuracy: 0.8792 - val_loss: 0.3282
```

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
```

```
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()
```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_labels = (predictions > 0.5).astype(int).flatten()
```

79/79 ————— 7s 90ms/step

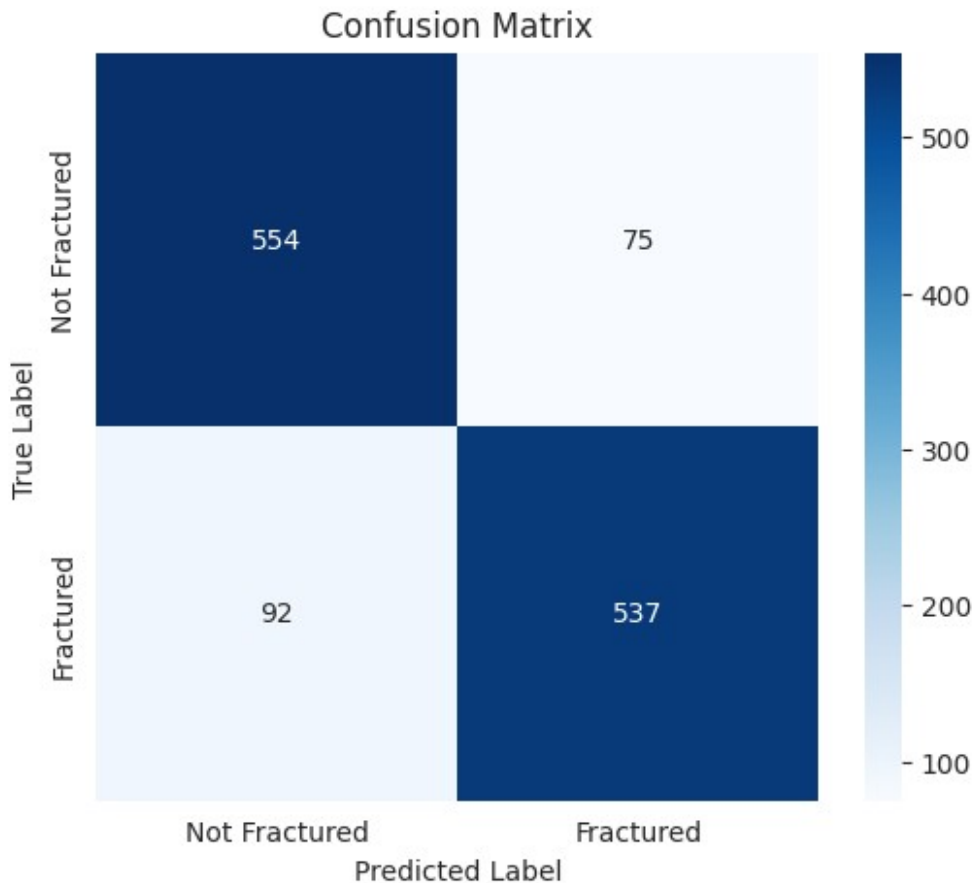
```
from sklearn.metrics import classification_report
```

```
report = classification_report(test_labels, predicted_labels,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
Benign	0.86	0.88	0.87	629
Malignant	0.88	0.85	0.87	629
accuracy			0.87	1258
macro avg	0.87	0.87	0.87	1258
weighted avg	0.87	0.87	0.87	1258

```
conf_matrix = confusion_matrix(test_labels, predicted_labels)
```

```
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Malignant', 'Benign'], yticklabels=['Malignant',
            'Benign'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()
```



```
from tensorflow.keras import models, layers, regularizers

def create_cnn_model(input_shape):
    model = models.Sequential()

    model.add(layers.Conv2D(32, (3, 3), activation='relu',
input_shape=input_shape,
kernel_regularizer=regularizers.l2(0.001)))
    model.add(layers.MaxPooling2D((2, 2)))
    model.add(layers.Dropout(0.25))

    model.add(layers.Conv2D(64, (3, 3), activation='relu',
```

```

kernel_regularizer=regularizers.l2(0.001)))
    model.add(layers.MaxPooling2D((2, 2)))
    model.add(layers.Dropout(0.25))

    model.add(layers.Conv2D(128, (3, 3), activation='relu',
kernel_regularizer=regularizers.l2(0.001)))
    model.add(layers.MaxPooling2D((2, 2)))
    model.add(layers.Dropout(0.25))

    model.add(layers.Flatten())

    model.add(layers.Dense(128, activation='relu',
                           kernel_regularizer=regularizers.l2(0.001)))

    model.add(layers.Dropout(0.5))
    model.add(layers.Dense(1, activation='sigmoid'))

    return model

input_shape = (224, 224, 3)
num_classes = 1

cnn_model = create_cnn_model(input_shape)

cnn_model.compile(optimizer='adam',
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=10,
    callbacks=[early_stopping],
    verbose=1
)

Epoch 1/10
629/629 ————— 36s 43ms/step - accuracy: 0.6592 - loss:
1.1970 - val_accuracy: 0.7965 - val_loss: 0.6146
Epoch 2/10
629/629 ————— 20s 31ms/step - accuracy: 0.8189 - loss:
0.5164 - val_accuracy: 0.8243 - val_loss: 0.5465
Epoch 3/10
629/629 ————— 21s 33ms/step - accuracy: 0.8201 - loss:
0.4959 - val_accuracy: 0.7909 - val_loss: 0.5115
Epoch 4/10
629/629 ————— 19s 30ms/step - accuracy: 0.8231 - loss:
0.4594 - val_accuracy: 0.8450 - val_loss: 0.4743


```


Epoch 5/10

629/629 ————— 20s 32ms/step - accuracy: 0.8371 - loss: 0.4437 - val\_accuracy: 0.8362 - val\_loss: 0.4590

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_labels = (predictions > 0.5).astype(int).flatten()
```

79/79 ————— 2s 28ms/step

```
report = classification_report(test_labels, predicted_labels,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
Benign	0.91	0.63	0.75	629
Malignant	0.72	0.94	0.81	629
accuracy			0.79	1258
macro avg	0.82	0.79	0.78	1258
weighted avg	0.82	0.79	0.78	1258

```
conf_matrix = confusion_matrix(test_labels, predicted_labels)
```


```
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
```



```

xticklabels=['Malignant', 'Benign'], yticklabels=['Malignant',
'Benign'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()

```



```

from tensorflow.keras.applications import Xception
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise
from tensorflow.keras.optimizers import Adam
import tensorflow as tf

def create_xception_model(input_shape):

    base_model = Xception(weights='imagenet', input_shape=input_shape,
include_top=False)

    for layer in base_model.layers:
        layer.trainable = False

```

```

model = Sequential()
model.add(base_model)
model.add(GaussianNoise(0.25))
model.add(GlobalAveragePooling2D())
model.add(Dense(512, activation='relu'))
model.add(BatchNormalization())
model.add(GaussianNoise(0.25))
model.add(Dropout(0.25))
model.add(Dense(1, activation='sigmoid'))

```

```

return model

```

```

input_shape = (224, 224, 3)
cnn_model = create_xception_model(input_shape)

```

```

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/xception/xception\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/xception/xception_weights_tf_dim_ordering_tf_kernels_notop.h5)  
83683744/83683744 ————— 0s 0us/step

```

history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=10,
    callbacks=[early_stopping],
    verbose=1
)

```

Epoch 1/10

629/629 ————— 78s 97ms/step - accuracy: 0.7860 - loss: 0.4897 - val\_accuracy: 0.8641 - val\_loss: 0.3278

Epoch 2/10

629/629 ————— 42s 67ms/step - accuracy: 0.8590 - loss: 0.3347 - val\_accuracy: 0.8768 - val\_loss: 0.3074

Epoch 3/10

629/629 ————— 42s 66ms/step - accuracy: 0.8758 - loss: 0.2945 - val\_accuracy: 0.8760 - val\_loss: 0.3123

Epoch 4/10

629/629 ————— 83s 67ms/step - accuracy: 0.8856 - loss: 0.2739 - val\_accuracy: 0.8752 - val\_loss: 0.2933

Epoch 5/10

629/629 ————— 82s 67ms/step - accuracy: 0.8927 - loss: 0.2578 - val\_accuracy: 0.8831 - val\_loss: 0.2869


Epoch 6/10


629/629 ————— 42s 67ms/step - accuracy: 0.8926 - loss:

```
0.2599 - val_accuracy: 0.8847 - val_loss: 0.2827
Epoch 7/10
629/629 _____ 42s 66ms/step - accuracy: 0.9089 - loss:
0.2282 - val_accuracy: 0.8808 - val_loss: 0.2843
Epoch 8/10
629/629 _____ 42s 67ms/step - accuracy: 0.9093 - loss:
0.2160 - val_accuracy: 0.8696 - val_loss: 0.3361
Epoch 9/10
629/629 _____ 42s 66ms/step - accuracy: 0.9194 - loss:
0.2026 - val_accuracy: 0.8903 - val_loss: 0.2875
Epoch 10/10
629/629 _____ 42s 67ms/step - accuracy: 0.9204 - loss:
0.1967 - val_accuracy: 0.8927 - val_loss: 0.2766
```

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```





```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_labels = (predictions > 0.5).astype(int).flatten()
```

79/79 ————— 9s 87ms/step

```
report = classification_report(test_labels, predicted_labels,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
Benign	0.86	0.93	0.89	629
Malignant	0.92	0.85	0.88	629
accuracy			0.89	1258
macro avg	0.89	0.89	0.89	1258
weighted avg	0.89	0.89	0.89	1258


```
conf_matrix = confusion_matrix(test_labels, predicted_labels)
```

```
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
```

```

xticklabels=['Malignant', 'Benign'], yticklabels=['Malignant',
'Benign'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()

```



```

from tensorflow.keras.applications import InceptionV3
def create_inception_model(input_shape):
    base_model = InceptionV3(weights='imagenet',
input_shape=input_shape, include_top=False)

    for layer in base_model.layers:
        layer.trainable = False

    model = Sequential()
    model.add(base_model)

    model.add(GaussianNoise(0.25))

```

```

model.add(GlobalAveragePooling2D())

model.add(Dense(512, activation='relu'))
model.add(BatchNormalization())
model.add(GaussianNoise(0.25))
model.add(Dropout(0.25))

model.add(Dense(1, activation='sigmoid'))

return model

from tensorflow.keras.layers import GlobalAveragePooling2D, Dense,
Dropout, BatchNormalization, GaussianNoise

input_shape = (224, 224, 3)
cnn_model = create_inception_model(input_shape)

cnn_model.compile(optimizer=Adam(learning_rate=0.0001),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

early_stopping = EarlyStopping(monitor='val_loss', patience=5,
restore_best_weights=True)

history = cnn_model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=10,
    callbacks=[early_stopping],
    verbose=1
)

```

Epoch 1/10

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1729686928.096262 139 service.cc:145] XLA service 0x7aaaae00156e0 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

I0000 00:00:1729686928.096350 139 service.cc:153] StreamExecutor device (0): Tesla T4, Compute Capability 7.5

I0000 00:00:1729686928.096359 139 service.cc:153] StreamExecutor device (1): Tesla T4, Compute Capability 7.5

2/629 ————— 40s 64ms/step - accuracy: 0.4219 - loss: 0.9532

I0000 00:00:1729686940.121983 139 device\_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.


```
629/629 _____ 97s 121ms/step - accuracy: 0.7807 - loss:
0.4874 - val_accuracy: 0.8561 - val_loss: 0.3548
Epoch 2/10
629/629 _____ 26s 40ms/step - accuracy: 0.8461 - loss:
0.3580 - val_accuracy: 0.8672 - val_loss: 0.3203
Epoch 3/10
629/629 _____ 26s 42ms/step - accuracy: 0.8582 - loss:
0.3326 - val_accuracy: 0.8768 - val_loss: 0.3135
Epoch 4/10
629/629 _____ 30s 46ms/step - accuracy: 0.8694 - loss:
0.3078 - val_accuracy: 0.8665 - val_loss: 0.3426
Epoch 5/10
629/629 _____ 26s 40ms/step - accuracy: 0.8782 - loss:
0.2898 - val_accuracy: 0.8704 - val_loss: 0.3263
Epoch 6/10
629/629 _____ 26s 41ms/step - accuracy: 0.8821 - loss:
0.2773 - val_accuracy: 0.8712 - val_loss: 0.3087
Epoch 7/10
629/629 _____ 26s 40ms/step - accuracy: 0.8899 - loss:
0.2671 - val_accuracy: 0.8577 - val_loss: 0.3121
Epoch 8/10
629/629 _____ 26s 41ms/step - accuracy: 0.8959 - loss:
0.2478 - val_accuracy: 0.8760 - val_loss: 0.3033
Epoch 9/10
629/629 _____ 25s 40ms/step - accuracy: 0.8993 - loss:
0.2471 - val_accuracy: 0.8665 - val_loss: 0.3149
Epoch 10/10
629/629 _____ 25s 40ms/step - accuracy: 0.9114 - loss:
0.2367 - val_accuracy: 0.8641 - val_loss: 0.3243
```

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```







```
test_labels = test_gen_new.classes
predictions = cnn_model.predict(test_gen_new)
predicted_labels = (predictions > 0.5).astype(int).flatten()
```

79/79 ————— 16s 146ms/step

```
report = classification_report(test_labels, predicted_labels,
target_names=list(test_gen_new.class_indices.keys()))
print(report)
```

	precision	recall	f1-score	support
Benign	0.90	0.87	0.89	629
Malignant	0.88	0.90	0.89	629
accuracy			0.89	1258
macro avg	0.89	0.89	0.89	1258
weighted avg	0.89	0.89	0.89	1258


```
conf_matrix = confusion_matrix(test_labels, predicted_labels)
```

```
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
```

```

xticklabels=['Malignant', 'Benign'], yticklabels=['Malignant',
'Benign'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()

```



```

from tensorflow.keras.preprocessing.image import load_img,
img_to_array

def preprocess_image(image_path, target_size=(224, 224)):
    img = load_img(image_path, target_size=target_size)
    img_array = img_to_array(img)
    img_array = np.expand_dims(img_array, axis=0)
    img_array = img_array / 255.0
    return img_array

def predict_image(model, image_path):
    processed_img = preprocess_image(image_path)
    prediction = model.predict(processed_img)
    return prediction

```

```

image_path =
'/kaggle/input/melanoma-cancer-dataset/test/Malignant/5602.jpg'
prediction = predict_image(cnn_model, image_path)

if prediction[0][0] >= 0.5:
    print(f"Prediction: Malignant ({prediction[0][0]:.2f}
confidence)")
else:
    print(f"Prediction: Benign ({prediction[0][0]:.2f} confidence)")

1/1 ————— 0s 23ms/step
Prediction: Malignant (0.81 confidence)

import time

def predict_image_with_time(model, image_path):
    processed_img = preprocess_image(image_path)
    start_time = time.time()
    prediction = model.predict(processed_img)
    end_time = time.time()
    prediction_time = end_time - start_time
    return prediction, prediction_time

image_path =
'/kaggle/input/melanoma-cancer-dataset/test/Malignant/5602.jpg'
prediction, prediction_time = predict_image_with_time(cnn_model,
image_path)

if prediction[0][0] >= 0.5:
    print(f"Prediction: Malignant ({prediction[0][0]:.2f}
confidence)")
else:
    print(f"Prediction: Benign ({prediction[0][0]:.2f} confidence)")

print(f"Time taken for prediction: {prediction_time:.4f} seconds")

1/1 ————— 0s 25ms/step
Prediction: Malignant (0.81 confidence)
Time taken for prediction: 0.0777 seconds

```

## 1. CNN Model:

- **Benign:**
  - Precision: 0.86
  - Recall: 0.88
  - F1-score: 0.87
- **Malignant:**
  - Precision: 0.88
  - Recall: 0.85
  - F1-score: 0.87

- **Overall Accuracy:** 0.87
- **Macro Avg/Weighted Avg F1-score:** 0.87

**Analysis:** The CNN model without regularization performs well with a balanced precision, recall, and f1-score for both classes. Overall, the accuracy is 87%, and the performance is consistent across metrics.

## 2. CNN Model with Regularization:

- **Benign:**
  - Precision: 0.91
  - Recall: 0.63
  - F1-score: 0.75
- **Malignant:**
  - Precision: 0.72
  - Recall: 0.94
  - F1-score: 0.81
- **Overall Accuracy:** 0.79
- **Macro Avg/Weighted Avg F1-score:** 0.78

**Analysis:** Adding regularization negatively impacted performance, especially for the **Benign** class, with a sharp decline in recall (0.63). The **Malignant** class benefits from improved recall but at the cost of reduced precision. Overall, accuracy and f1-scores dropped to 79%, indicating a potential imbalance in how the model is generalizing across classes.

## 3. Xception Model:

- **Benign:**
  - Precision: 0.86
  - Recall: 0.93
  - F1-score: 0.89
- **Malignant:**
  - Precision: 0.92
  - Recall: 0.85
  - F1-score: 0.88
- **Overall Accuracy:** 0.89
- **Macro Avg/Weighted Avg F1-score:** 0.89

**Analysis:** The Xception model shows strong performance with high precision and recall for both **Benign** and **Malignant** classes. Accuracy of 89% and balanced f1-scores indicate good generalization. This is a clear improvement over both the CNN and the regularized CNN models.

## 4. Inception Model:

- **Benign:**
  - Precision: 0.90
  - Recall: 0.87
  - F1-score: 0.89
- **Malignant:**

- Precision: 0.88
- Recall: 0.90
- F1-score: 0.89
- **Overall Accuracy: 0.89**
- **Macro Avg/Weighted Avg F1-score: 0.89**

**Analysis:** The Inception model performs similarly to Xception, with 89% accuracy and nearly equal precision, recall, and f1-scores across both classes. This model strikes a good balance between both **Benign** and **Malignant** detection.

## Conclusion:

- The **CNN with regularization** struggled with class imbalance, especially for the **Benign** class.
- **Xception and Inception** models achieved the best performance, with both reaching 89% accuracy and well-balanced precision, recall, and f1-scores.
- Between **Xception and Inception**, the performance is very close, and either model could be chosen based on other factors like computational efficiency or ease of deployment.

```
models = ['CNN', 'CNN + Reg', 'Xception', 'Inception']
precision = [0.87, 0.82, 0.89, 0.89]
recall = [0.87, 0.79, 0.89, 0.89]
f1_score = [0.87, 0.78, 0.89, 0.89]
accuracy = [0.87, 0.79, 0.89, 0.89]

bar_width = 0.2

r1 = np.arange(len(models))
r2 = [x + bar_width for x in r1]
r3 = [x + bar_width for x in r2]
r4 = [x + bar_width for x in r3]

plt.figure(figsize=(10, 6))
plt.bar(r1, precision, color='b', width=bar_width, edgecolor='grey',
label='Precision')
plt.bar(r2, recall, color='g', width=bar_width, edgecolor='grey',
label='Recall')
plt.bar(r3, f1_score, color='r', width=bar_width, edgecolor='grey',
label='F1-Score')
plt.bar(r4, accuracy, color='c', width=bar_width, edgecolor='grey',
label='Accuracy')

plt.xlabel('Models', fontweight='bold')
plt.ylabel('Scores', fontweight='bold')
plt.title('Comparative Performance of Models', fontweight='bold')
plt.xticks([r + bar_width for r in range(len(models))], models)

plt.legend()

plt.show()
```

